

AIT Semantic and Declarative Technologies Course
Assignment 1: Checking the consistency of a Sudoku puzzle
Deadline: Friday, October 16, 2015, 23:59

Description

This task is related to the well-known Sudoku puzzle.

A *Sudoku grid* is square matrix consisting of m rows and m columns where m itself is a square number, $m = k * k$. We refer to the elements of the matrix as the *fields* of the grid. The Sudoku grid subdivides to square *sub-grids* of k rows and k columns. For the mathematically minded, the top left fields of sub-grids have (row,column) coordinates $(i * k, j * k)$, for $i = 0, \dots, k - 1$ and $j = 0, \dots, k - 1$; assuming that the top left field of the whole grid has (row,column) coordinates $(0, 0)$. An example grid for $k = 2$ is shown below, the sub-grid borders are indicated by double lines.

2	3	4	1
4		2	
	2	1	4
	4		2

The task is to check the consistency of a partially filled Sudoku grid. The Sudoku grid is given as a list of rows, each row being a list of integers. You can assume (i.e. you don't have to check) that the Sudoku grid supplied to your program is correct in the following sense: it consists of m rows, each row is a list of m fields, and each field is an integer between 0 and m , where $m = k * k$. You can also assume that $1 \leq k \leq 10$, although probably you will not make use of this assumption.

The above example is supplied to your program as the following Prolog data structure:

```
[[2,3, 4,1],
 [4,0, 2,0],

 [0,2, 1,4],
 [0,4, 0,2]]
```

The integer 0 represents a field which is not yet filled in. Positive integers represent fields that are filled in.

A Sudoku grid is said to be *inconsistent* if it contains a row or a column or a sub-grid, which, in turn, contains two occurrences of the same positive integer. A Sudoku grid is said to be *consistent* if it is not inconsistent.

The name *area* is used to denote any of the rows, columns or sub-grids. Using this terminology, we can state that a Sudoku grid is consistent if no area contains two occurrences of the same positive integer.

The above example grid is consistent. The following example grid is inconsistent, because the bottom-right sub-grid contains two occurrences of integer 1.

```
[[2,3, 0,0],
 [1,0, 2,3],

 [3,2, 1,4],
 [4,0, 0,1]]
```

The assignment

Write a Prolog predicate `consistent/1` which takes a Sudoku grid and succeeds if, and only if, the grid is consistent. If the grid is consistent, the predicate should succeed exactly once.

A Sudoku grid is represented by a Prolog term which is a list whose elements represent the rows of the grid in top-down order. Each row is represented by a list of integers, the fields in the given row, in left-to-right order.

The predicate has the following specification (head comment):

```
% consistent(+SGrid): For all areas of the Sudoku grid SGrid it holds
% that all positive integers in the area are distinct.
```

Sample runs

```
| ?- consistent(  
    [[1]]  
    ).  
yes  
| ?- consistent(  
    [[2,3, 4,1],  
     [4,0, 2,0],  
     [0,2, 1,4],  
     [0,4, 0,2]]  
    ).  
yes  
| ?- consistent(  
    [[2,3, 0,0],  
     [1,0, 2,3],  
     [3,2, 1,4],  
     [4,0, 0,1]]  
    ).  
no  
| ?- consistent(  
    [[0,0, 0,0],  
     [0,0, 0,0],  
     [0,0, 0,0],  
     [0,0, 0,0]]  
    ).  
yes
```

Notes

- The following example shows how obtain the integer square root K of a square number M using the built-in predicate `is/2`:

```
K is integer(sqrt(M))
```
- Instructions for submitting solutions and downloading test cases will be announced soon.
- You may find useful some predicates in `library(lists)`, e.g. `transpose/2` and `sublist/5`. In solving this assignment you can freely use predicates from this library. Ask the instructor if you would like to use other libraries.

Optional extension

A Sudoku grid is said to be a *refinement* of another Sudoku grid, if the former can be obtained by *filling in* (zero or more) fields of value 0 in the latter, i.e. replacing the 0 value by a positive integer.

A Sudoku grid is said to be *fully filled in*, if all fields contain positive integers.

A Sudoku grid is said to be *complete*, if it is fully filled in and consistent.

As an optional extension you can implement the following predicate for solving a Sudoku puzzle:

```
% sudoku(+Grid0, ?Grid): Grid is a complete refinement of the Sudoku grid Grid0  
  
| ?- sudoku([[2,3, 4,1],[4,0, 2,0],[0,2, 1,4],[0,4, 0,2]], Grid).  
Grid = [[2,3,4,1],[4,1,2,3],[3,2,1,4],[1,4,3,2]] ? ;  
no  
| ?-
```

You can provide simple or more complex implementations for the above predicate. In increasing order of ambition and efficiency you can consider the following approaches:

1. **Generate-and-test:** Generate an arbitrary filled in refinement of the input grid and check if it happens to be consistent.
2. **Fused generate-and-test:** Refine the input grid by filling in a single field, and immediately check the result for consistency. Continue this process until a (fully) filled in refinement is obtained.
3. **Fill in fields that are unique:** Try to find empty fields that can only be assigned a single positive value while keeping the grid consistent.
4. **Find integers that can uniquely placed in an area:** Try to find integers that can only be placed into a single field of an area while keeping the grid consistent.
5. **Keep the braching factor of the search tree low:** If you run out of deterministic refinements (such as listed as the two preceding items), try to find a field assigning which leads to the smallest number of consistent refinements.